

# qhbmndohku

March 28, 2025

```
[3]: import pandas as pd
import numpy as np
import matplotlib as plt
```

```
[4]: df=pd.read_csv("social_network_ads.csv")
```

```
[5]: df
```

```
[5]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

```
[6]: df.columns
```

```
[6]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'],
dtype='object')
```

```
[7]: df.isnull()
```

```
[7]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
..	...	...	...	...	...

```

395     False     False     False           False     False
396     False     False     False           False     False
397     False     False     False           False     False
398     False     False     False           False     False
399     False     False     False           False     False

```

[400 rows x 5 columns]

```

[8]: from sklearn.preprocessing import LabelEncoder
     le = LabelEncoder()
     df['Gender'] = le.fit_transform(df['Gender'])
     newdf=df

```

```

[9]: df

```

```

[9]:      User ID  Gender  Age  EstimatedSalary  Purchased
0      15624510        1   19             19000           0
1      15810944        1   35             20000           0
2      15668575        0   26             43000           0
3      15603246        0   27             57000           0
4      15804002        1   19             76000           0
..      ...      ...  ...      ...      ...
395     15691863        0   46             41000           1
396     15706071        1   51             23000           1
397     15654296        0   50             20000           1
398     15755018        1   36             33000           0
399     15594041        0   49             36000           1

```

[400 rows x 5 columns]

```

[29]: # Splitting the DataFrame into features (X) and target variable (y)
     x = df.drop(['Purchased'], axis=1) # X: Features (all columns except
     ↪ 'Purchased')
     y = df['Purchased']                # y: Target variable (Purchased status)

```

```

[30]: # Importing the train_test_split function to split the dataset into training
     ↪ and testing sets
     from sklearn.model_selection import train_test_split

```

```

[31]: # Splitting the dataset into training and testing sets
     # X_train, Y_train → Training data (60%)
     # X_test, Y_test → Testing data (40%)
     X_train, X_test, Y_train, Y_test = train_test_split(x, y, test_size=0.4,
     ↪ random_state=10)

```

```

[32]: # Importing the LogisticRegression model from sklearn for binary classification
     from sklearn.linear_model import LogisticRegression

```

```
[34]: # Displaying the first 5 rows of the training feature set
print(X_train.head())
```

```
      Gender
0      Male
4      Male
1    Female
```

```
[36]: # One-hot encoding for categorical variables
# drop_first=True → Avoids multicollinearity by dropping the first category
X_train = pd.get_dummies(X_train, drop_first=True)
X_test = pd.get_dummies(X_test, drop_first=True)

# Initializing the Logistic Regression model
logreg = LogisticRegression()

# Fitting the model to the training data
logreg.fit(X_train, Y_train)
```

```
[36]: LogisticRegression()
```

```
[37]: # Making predictions on the testing set using the trained logistic regression
      ↪ model
Y_pred = logreg.predict(X_test) # Predicting the target values for X_test

# Printing the predicted values
print("Predictions:", Y_pred)
```

```
Predictions: ['No' 'No']
```

```
[38]: # Importing the sklearn library and LogisticRegression model
import sklearn
from sklearn.linear_model import LogisticRegression # Logistic regression model

# Initializing the Logistic Regression model
logreg = LogisticRegression()

# Fitting the model to the training data
model = logreg.fit(X_train, Y_train) # model contains the trained logistic
      ↪ regression instance
```

```
[39]: # Making predictions on the training set
Ytrain_pred = logreg.predict(X_train) # Predicted labels for the training set

# Making predictions on the testing set
Ytest_pred = logreg.predict(X_test) # Predicted labels for the testing set
```

```
[40]: import pandas as pd

# Creating DataFrame for training predictions
df_train = pd.DataFrame({'Actual': Y_train, 'Predicted': Ytrain_pred})

# Creating DataFrame for testing predictions
df_test = pd.DataFrame({'Actual': Y_test, 'Predicted': Ytest_pred})

# Displaying the first few rows
print("Training Predictions:")
print(df_train.head())

print("\nTesting Predictions:")
print(df_test.head())
```

Training Predictions:

	Actual	Predicted
0	Yes	No
4	No	No
1	No	No

Testing Predictions:

	Actual	Predicted
2	Yes	No
3	Yes	No

```
[41]: # Importing evaluation metrics from sklearn
from sklearn.metrics import precision_score, confusion_matrix, accuracy_score, recall_score

# Creating the confusion matrix for the testing set
cm = confusion_matrix(Y_test, Y_pred) # Compares actual vs predicted labels
```

```
[42]: # Creating the confusion matrix for the testing set
cm = confusion_matrix(Y_test, Y_pred) # Compares actual vs predicted labels
```

```
[43]: # Generating the confusion matrix for the testing set
cm = confusion_matrix(Y_test, Y_pred) # Compares actual vs predicted labels

# Printing the confusion matrix
print("Confusion Matrix:\n", cm)
```

Confusion Matrix:

```
[[0 0]
 [2 0]]
```

```
[45]: # Printing the model's performance metrics
print("Accuracy:", accuracy_score(Y_test, Y_pred))           # Overall
    ↳ accuracy of the model
print("Precision:", precision_score(Y_test, Y_pred, average='weighted')) #
    ↳ Precision considering class imbalance
print("Recall:", recall_score(Y_test, Y_pred, average='weighted'))   #
    ↳ Recall considering class imbalance
```

Accuracy: 0.0

Precision: 0.0

Recall: 0.0

C:\Users\Vishwajeet Kulkarni\anaconda3\Lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))  
C:\Users\Vishwajeet Kulkarni\anaconda3\Lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

```
[48]: # Importing necessary libraries
import pandas as pd           # For creating and handling DataFrames
import matplotlib.pyplot as plt # For plotting visualizations

# Sample DataFrame (replace with your actual data)
data = {
    'Gender': ['Male', 'Female', 'Male', 'Female', 'Male'], # Gender column
    'Purchased': ['Yes', 'No', 'Yes', 'Yes', 'No']           # Purchased status
}

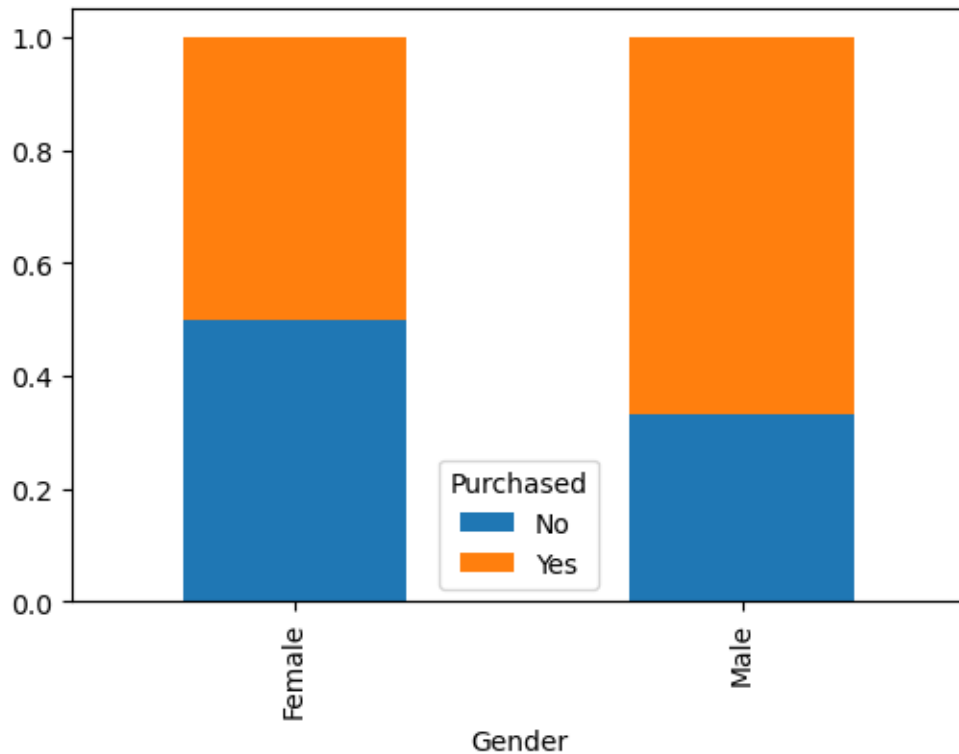
# Creating a DataFrame from the sample data
df = pd.DataFrame(data)

# Creating a crosstab to show the proportion of purchases by gender
# normalize="index" → Normalizes the values by row (percentage distribution)
cm = pd.crosstab(df['Gender'], df['Purchased'], normalize="index")
print(cm) # Displaying the normalized crosstab

# Plotting the crosstab as a stacked bar chart
cm.plot.bar(figsize=(6, 4), stacked=True) # Stacked bar chart with dimensions
    ↳ 6x4
plt.show() # Displaying the plot
```

	Purchased	No	Yes
Gender			

Female	0.500000	0.500000
Male	0.333333	0.666667




```
[46]: # Importing Seaborn for data visualization
import seaborn as sns

# Plotting the distribution of the 'Purchased' column
sns.displot(df['Purchased']) # Distribution plot of the target variable
```

C:\Users\Vishwajeet Kulkarni\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

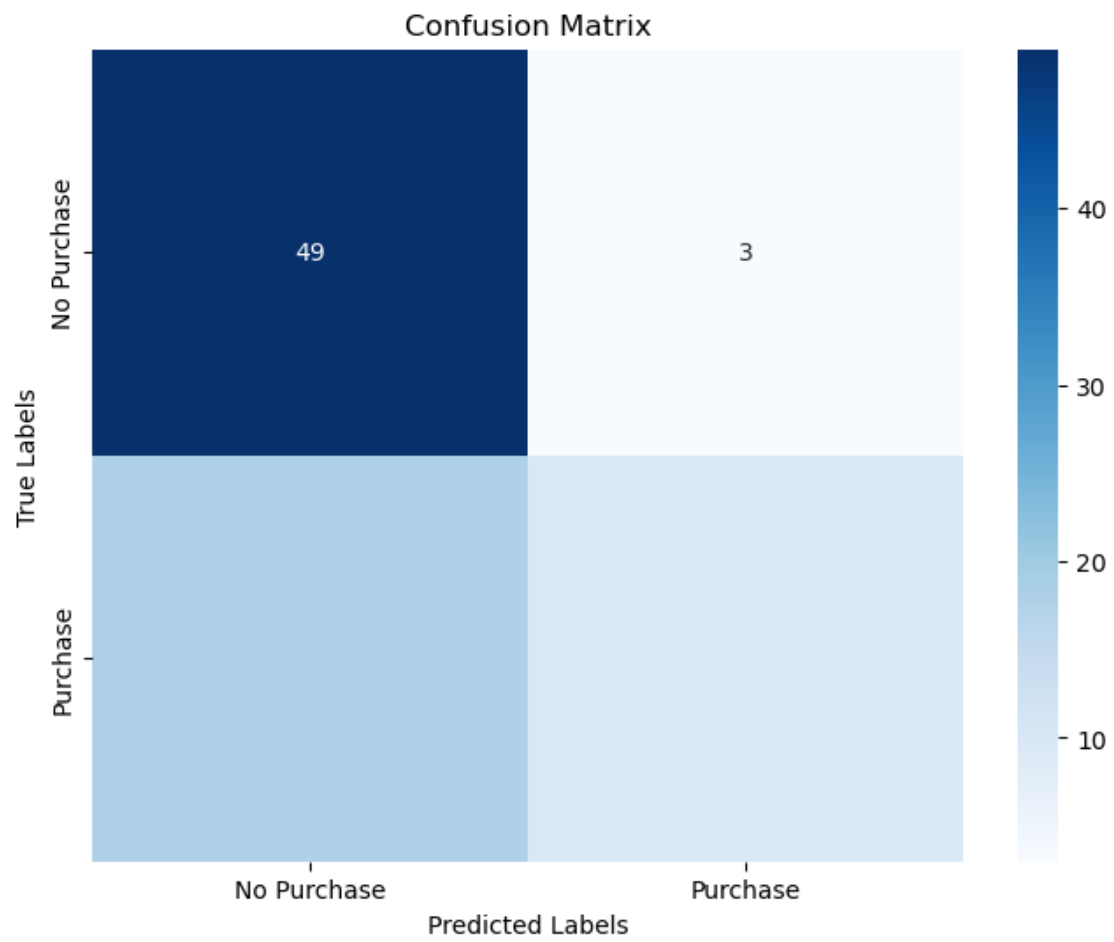
```
[46]: <seaborn.axisgrid.FacetGrid at 0x18038f5e390>
```



```
[47]: # Importing necessary libraries for visualization and evaluation
import matplotlib.pyplot as plt      # For creating visual plots
import seaborn as sns                # For enhanced visualizations
from sklearn.metrics import confusion_matrix # For generating the confusion matrix
```

```
[42]: # Generate confusion matrix
cm = confusion_matrix(Y_test, Y_pred)

# Plot confusion matrix with correct labels
labels = ['No Purchase', 'Purchase'] # Your class labels
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels,
            yticklabels=labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```



[ ]: