

zinigvvch

March 28, 2025

```
[1]: # Importing necessary libraries
import pandas as pd          # For creating and handling DataFrames
import numpy as np          # For numerical operations and array
    ↪ handling
import matplotlib as plt    # For data visualization (matplotlib
    ↪ module)
```

```
[2]: # Loading the dataset into a DataFrame from a CSV file
df = pd.read_csv("diabetes.csv") # Reads the CSV file and stores it in the
    ↪ DataFrame 'df'
```

```
[3]: df
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1

767                      0.315    23            0

[768 rows x 9 columns]

```
[4]: df.columns
```

```
[4]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[5]: df.isnull()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
..	...	...	...	...	...	...	
763	False	False	False	False	False	False	
764	False	False	False	False	False	False	
765	False	False	False	False	False	False	
766	False	False	False	False	False	False	
767	False	False	False	False	False	False	

	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
..	...	...	...
763	False	False	False
764	False	False	False
765	False	False	False
766	False	False	False
767	False	False	False

[768 rows x 9 columns]

```
[3]: # Splitting the dataset into features (X) and target variable (y)

# X → Features (all columns except 'Outcome')
x = df.drop(['Outcome'], axis=1)

# y → Target variable (Outcome column)
y = df['Outcome']
```

```
[4]: # Importing the train_test_split function from sklearn for data splitting
from sklearn.model_selection import train_test_split

# Splitting the dataset into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(x,y,test_size=0.
↪4,random_state=10)
```

```
[5]: # Importing the Gaussian Naive Bayes classifier from sklearn
from sklearn.naive_bayes import GaussianNB

# Initializing the Gaussian Naive Bayes model
gaussian = GaussianNB()

# Fitting the model to the training data
gaussian.fit(X_train, Y_train) # Trains the model using the training set
```

```
[5]: GaussianNB()
```

```
[6]: # Making predictions on the testing set using the trained Naive Bayes model
Y_pred = gaussian.predict(X_test) # Predicted labels for the testing set
```

```
[7]: # Importing performance evaluation metrics from sklearn
from sklearn.metrics import accuracy_score, precision_score, recall_score
```

```
[8]: # Calculating model performance metrics

# Accuracy: Overall correctness of the model
accuracy = accuracy_score(Y_test, Y_pred)

# Precision: Proportion of true positives among all positive predictions
precision = precision_score(Y_test, Y_pred, average='micro')

# Recall: Proportion of true positives identified correctly
recall = recall_score(Y_test, Y_pred, average='micro')
```

```
[9]: # Importing performance evaluation metrics from sklearn
from sklearn.metrics import precision_score, confusion_matrix, accuracy_score, ↵
↪recall_score

# Generating the confusion matrix for the testing set
cm = confusion_matrix(Y_test, Y_pred) # Compares actual vs predicted labels
```

```
[20]: cm = confusion_matrix(Y_test, Y_pred)
print("ConfusionMatrix:\n", cm)
```

```
ConfusionMatrix:
[[166  35]
```

```
[ 47  60]]
```

```
[21]: print("Accuracy:", accuracy_score(Y_test, Y_pred))
      print("Precision:", precision_score(Y_test, Y_pred, average='weighted'))
      print("Recall:", recall_score(Y_test, Y_pred, average='weighted'))
```

```
Accuracy: 0.7337662337662337
Precision: 0.7280092035466387
Recall: 0.7337662337662337
```

```
[10]: Y_pred
```

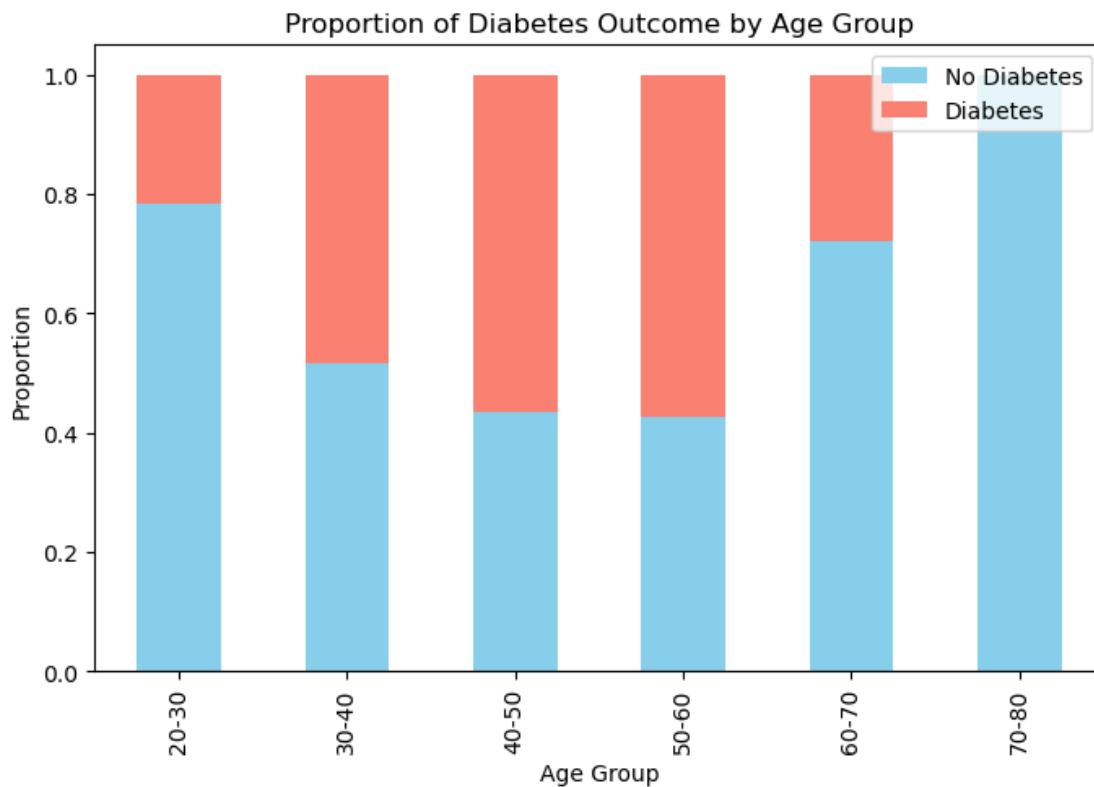
```
[10]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1,
          1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
          1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
          1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1,
          1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1,
          1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0,
          1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0,
          1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0,
          0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0,
          0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1,
          0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
          0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
          1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
          0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0],
          dtype=int64)
```

```
[13]: import matplotlib.pyplot as plt
      # Creating a crosstab to show the proportion of diabetes outcome by age groups
      # Binning the 'Age' column into categories for better visualization
      df['AgeGroup'] = pd.cut(df['Age'], bins=[20, 30, 40, 50, 60, 70, 80],
                              labels=['20-30', '30-40', '40-50', '50-60', '60-70', '70-80'])
      # Crosstab showing the proportion of diabetes outcome by age group
      cm = pd.crosstab(df['AgeGroup'], df['Outcome'], normalize="index")
      print("\nCrosstab of Age Group vs Outcome:")
      print(cm)

      # Plotting the crosstab as a stacked bar chart
      cm.plot.bar(figsize=(8, 5), stacked=True, color=['skyblue', 'salmon'])
      plt.title("Proportion of Diabetes Outcome by Age Group")
      plt.xlabel("Age Group")
      plt.ylabel("Proportion")
      plt.legend(['No Diabetes', 'Diabetes'])
      plt.show()
```

Crosstab of Age Group vs Outcome:

Outcome	0	1
AgeGroup		
20-30	0.784173	0.215827
30-40	0.515924	0.484076
40-50	0.433628	0.566372
50-60	0.425926	0.574074
60-70	0.720000	0.280000
70-80	1.000000	0.000000




```
[23]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
```

```
[24]: # Generate confusion matrix
cm = confusion_matrix(Y_test, Y_pred)

# Correct class labels
labels = ['No Diabetes', 'Diabetes'] # Use appropriate labels

# Plot confusion matrix
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels,
            yticklabels=labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```



[ ]: